

Package demo

William Shih, Ricardo Simpao, Nilay Varshney, Luke Yee

3/20/2020

Here is an example of how our package functions run. For our data set, we are using a “SGEMM GPU kernel performance Data Set,” which measures the running times of a matrix-matrix product, given different parameter combinations.

Below, all 4 functions in the package (calculating linear regression bootstrap, calculating confidence intervals for coefficients, prediction intervals, and confidence intervals for σ^2) are much faster using C++ than using R. Overhead with complex functions such as map, apply, and reduce may takes a large amount of time with R so that the C++ version that only uses RcppArmadillo and functions from std namespace is much faster. Use of syntactic sugar from Rcpp is minimized in the C++ functions.

```
library(devtools)
library(tidyverse)
library(STA141CFinal)
library(furrr)

set.seed(141)
dat = read_csv("sgemm_product.csv")
dat = dat[sample(241000, 1000),]
dat2 = dat[1:100,]

#We specify a specific column set
y = dat$`Run1 (ms)`
x = dat[,1:(ncol(dat)-4)]

#linear model objects
fit = linear_reg_bs_C(x, y, s = 10, r = 1000)
```

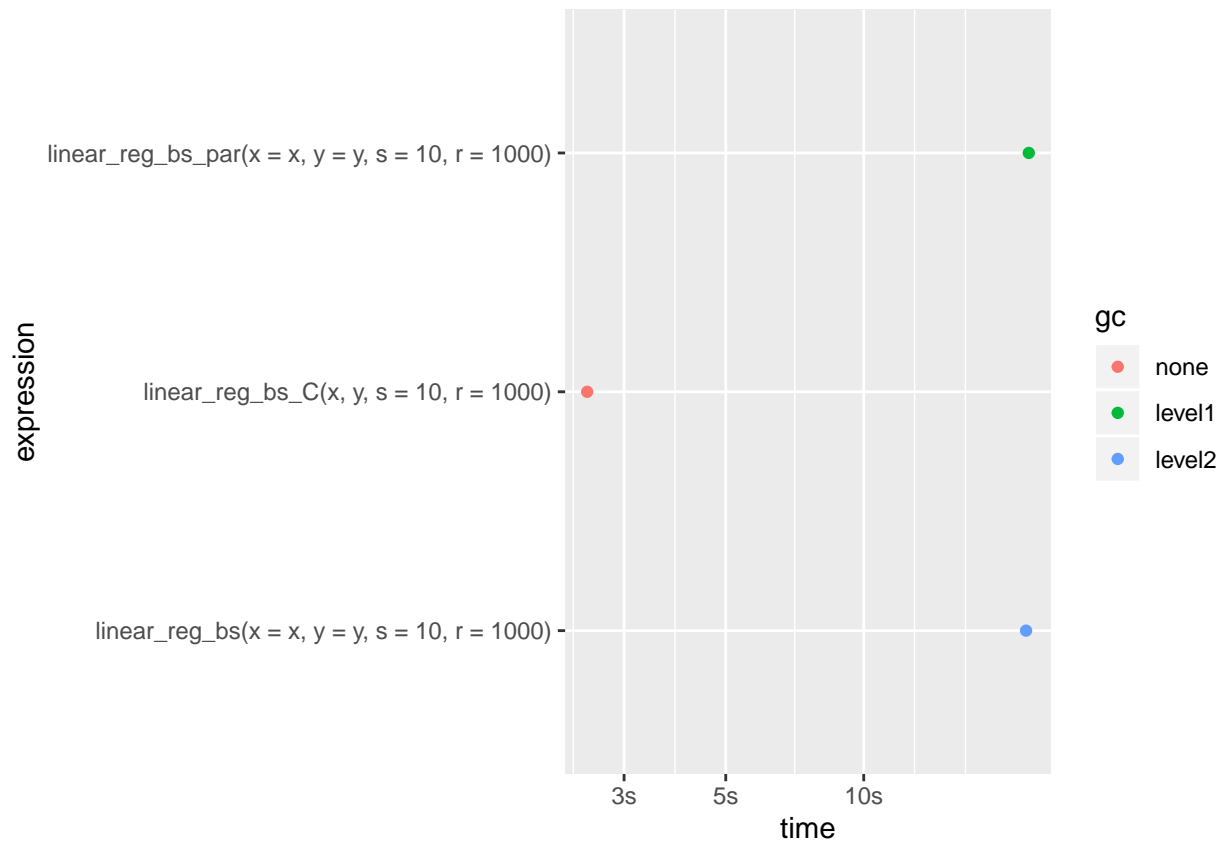
Linear Regression with blb (n = 1000, p = 15, subsets = 10, and replications = 1000)

```
## Warning: Some expressions had a GC in every iteration; so filtering is disabled.
```

```
## # A tibble: 3 x 6
##   expression                                min median `itr/sec`
##   <bch:expr>                                <bch:> <bch:>      <dbl>
## 1 linear_reg_bs(x = x, y = y, s = 10, r = 1000)    22.61s 22.61s    0.0442
## 2 linear_reg_bs_par(x = x, y = y, s = 10, r = 1000) 22.92s 22.92s    0.0436
## 3 linear_reg_bs_C(x, y, s = 10, r = 1000)         2.49s  2.49s    0.401
## # ... with 2 more variables: mem_alloc <bch:byt>, `gc/sec` <dbl>
```

```
## Warning in f(...): The default behavior of beeswarm has changed in version
## 0.6.0. In versions <0.6.0, this plot would have been dodged on the y-axis. In
## versions >=0.6.0, grouponX=FALSE must be explicitly set to group on y-axis.
```

```
## Please set grouponX=TRUE/FALSE to avoid this warning and ensure proper axis
## choice.
```



The C++ version is about 10 times faster than either of the R versions. The C++ version uses RcppArmadillo to multiply matrices as that was found to be the fastest version available. RcppArmadillo was generally faster than using multiplying matrices using `std::inner_product`. The R parallel version took just as long as the R non-parallel version.

95 % Confidence Interval for Variable Coefficients (original dataset has 1000 replications and 10 subsets)

```
coef_CI(fit, alpha = 0.05)
```

##	Lower_Bounds	Estimates	Upper_Bounds
## Intercept	-250.9593638	-161.822787	-77.756876
## MWG	2.4803281	2.909166	3.349526
## NWG	2.3378785	2.785394	3.245074
## KWG	4.6214054	6.870881	9.220991
## MDIMC	-18.6137757	-15.942214	-13.365062
## NDIMC	-18.0645940	-15.529066	-13.084430
## MDIMA	0.8616486	2.773963	4.808043
## NDIMB	-0.1020447	1.983522	4.069686
## KWI	0.8409772	5.744059	10.709685
## VWM	-1.5499689	9.779875	20.961781
## VWN	-4.8347615	4.594360	14.276281

```
## STRM      -11.0345567   19.774674   51.951853
## STRN      -57.5970022  -24.403050    9.056510
## SA         0.2486218   28.977774   58.002116
## SB         24.4859861   57.913604   92.355236
```

```
coef_CI_par(fit,alpha = 0.05)
```

```
##           Lower_Bounds   Estimates Upper_Bounds
## Intercept -250.9593638 -161.822787  -77.756876
## MWG         2.4803281   2.909166    3.349526
## NWG         2.3378785   2.785394    3.245074
## KWG         4.6214054   6.870881    9.220991
## MDIMC       -18.6137757 -15.942214  -13.365062
## NDIMC       -18.0645940 -15.529066  -13.084430
## MDIMA        0.8616486   2.773963    4.808043
## NDIMB       -0.1020447   1.983522    4.069686
## KWI         0.8409772   5.744059   10.709685
## VWM         -1.5499689   9.779875   20.961781
## VWN         -4.8347615   4.594360   14.276281
## STRM       -11.0345567   19.774674   51.951853
## STRN       -57.5970022  -24.403050    9.056510
## SA          0.2486218   28.977774   58.002116
## SB          24.4859861   57.913604   92.355236
```


```
coef_CI_C(fit,alpha = 0.05)
```

```
##           Lower_Bounds   Estimates Upper_Bounds
## Intercept -250.9291528 -161.822787  -76.566608
## MWG         2.4804299   2.909166    3.353727
## NWG         2.3379711   2.785394    3.249154
## KWG         4.6219064   6.870881    9.246523
## MDIMC       -18.6132246 -15.942214  -13.333145
## NDIMC       -18.0640532 -15.529066  -13.067617
## MDIMA        0.8620391   2.773963    4.834013
## NDIMB       -0.1017314   1.983522    4.088236
## KWI         0.8419045   5.744059   10.754005
## VWM         -1.5468762   9.779875   21.047596
## VWN         -4.8335146   4.594360   14.337845
## STRM       -11.0248991   19.774674   52.324871
## STRN       -57.5908069  -24.403050    9.346352
## SA          0.2563944   28.977774   58.204637
## SB          24.4928834   57.913604   92.451715
```

```
(b1 = bench::mark(
  coef_CI(fit, alpha = 0.05),
  coef_CI_par(fit,alpha = 0.05),
  coef_CI_C(fit, alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
##   expression          min    median `itr/sec` mem_alloc `gc/sec`
##   <bch:expr>      <bch:tm> <bch:tm>    <dbl>   <bch:byt>    <dbl>
## 1 coef_CI(fit, alpha = 0.05) 46.65ms    49ms    19.5    7.49MB    2.17
## 2 coef_CI_par(fit, alpha = 0.05) 83.27ms   106ms    9.75    7.8MB     0
## 3 coef_CI_C(fit, alpha = 0.05)  8.83ms    11ms   76.5    1.15MB     0
```

```
ggplot2::autoplot(b1)
```



The C++ version was only about 5 times faster than the R version. The C++ uses 1/7 as much memory as the R version. Also note that the C++ version calculates the quantiles differently from the R version. Therefore, the lower and upper bounds are slightly different in the C++ and R versions.

95% Prediction Interval (with $n = 100$ and $p = 14$ (original dataset has 1000 replications and 10 subsets))

```
plan(multiprocess, workers = 4)
PI(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5224  -64.28095    -12.19813
## [2,]     303.9097  373.82120     446.12019
## [3,]    -105.6873  -45.87250      11.50187
```

```
PI_par(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5224  -64.28095    -12.19813
## [2,]     303.9097  373.82120     446.12019
## [3,]    -105.6873  -45.87250      11.50187
```


```
PI_C(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5031  -64.28095    -11.66223
## [2,]     303.9193  373.82120    446.62104
## [3,]    -105.6698  -45.87250     12.00876
```

```
(b2 = bench::mark(
  PI(fit, dat2[1:3, 1:14], alpha = 0.05),
  PI_par(fit, dat2[1:3, 1:14], alpha = 0.05),
  PI_C(fit, dat2[1:3, 1:14], alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
##   expression          min  median `itr/sec`
##   <bch:expr>      <bch:tm> <bch:tm>    <dbl>
## 1 PI(fit, dat2[1:3, 1:14], alpha = 0.05)    54.68ms  65.18ms     14.4
## 2 PI_par(fit, dat2[1:3, 1:14], alpha = 0.05) 778.54ms 778.54ms     1.28
## 3 PI_C(fit, dat2[1:3, 1:14], alpha = 0.05)   2.35ms   2.91ms    340.
## # ... with 2 more variables: mem_alloc <bch:byt>, `gc/sec` <dbl>
```

```
ggplot2::autoplot(b2)
```



The C++ version to calculate the 100 95% confidence intervals was more than 20 times faster than the R version. Again, the C++ and R confidence intervals are slightly difference due to calculating the quantiles differently.

95 % Confidence Interval for Variance (original dataset has 1000 replications and 10 subsets)

```
s2_CI(fit, alpha = 0.05)
```

```
## Lower_Bound   Estimate Upper_Bound
##      572738.9    718571.7    858566.8
```

```
s2_CI_par(fit, alpha = 0.05)
```

```
## Lower_Bound   Estimate Upper_Bound
##      572738.9    718571.7    858566.8
```


```
s2_CI_C(fit, alpha = 0.05)
```

```
## Lower_Bound   Estimate Upper_Bound
##      572781.4    718571.7    859655.6
```

```
(b3 = bench::mark(
  s2_CI(fit, alpha = 0.05),
  s2_CI_par(fit, alpha = 0.05),
  s2_CI_C(fit, alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
##   expression                min    median `itr/sec` mem_alloc `gc/sec`
##   <bch:expr>          <bch:tm> <bch:tm>    <dbl>   <bch:byt>   <dbl>
## 1 s2_CI(fit, alpha = 0.05)    1.67ms  1.92ms    461.     118.12KB     0
## 2 s2_CI_par(fit, alpha = 0.05) 97.81ms 102.18ms    9.56      2.53MB    2.39
## 3 s2_CI_C(fit, alpha = 0.05)   97.4us  98.6us   8910.       2.49KB     0
```

```
ggplot2::autoplot(b3)
```



The C++ version is about 19 times faster than the R version. The parallel version is not optimal for this case as it takes 100 times longer than the non-parallel R version. For some reason, the parallel version is not found to be faster than the non-parallel version for any of the 4 functions. But the C++ function is much faster anyways, so that is the one to use when trying to calculate these as fast as possible.